

PM1: (avg. accuracy = 88.9%)

In PM1 we implement perceptron algorithm without normalising the data and dropping rows which possess NaN values. We have final results as shown below.

In perceptron, we find the coefficient values of a hyperplane which is our decision boundary. We do this by iterating through our data many times, and adding or subtracting a particular row based on whether our existing hyperplane (which was initialised to be 0 in the beginning) correctly or incorrec

PM1

AVERAGE ACCURACY 0.8887096774193548

AVERAGE PRECISION 0.9252874912501537

AVERAGE RECALL 0.8969063267086946

PM1

ACCURACY VARIANCE: 0.05156808089583569

PRECISION VARIANCE: 0.05388796638889157

RECALL VARIANCE: 0.09983908130766614

PM3: (avg. accuracy = 95.2%)

In PM3 we implement perceptron algorithm after normalising the data and we take mean of the respective features and replace NaNs. We have final results as shown below.

PM3

AVERAGE ACCURACY 0.9521276595744681

AVERAGE PRECISION 0.9556678661004643

AVERAGE RECALL 0.9674085171501403

PM3

ACCURACY VARIANCE: 0.017809222385939592

PRECISION VARIANCE: 0.007935132969455616

RECALL VARIANCE: 0.02924692172947626

PM4: (avg. accuracy = 95.2%)

In PM4 the only difference with respect to PM3 is that we shuffle columns of the dataset which clearly doesn't affect any of the measured parameters for perceptron. We have final results as shown below.

PM4

ACCURACY VARIANCE: 0.017809222385939592
PRECISION VARIANCE: 0.007935132969455616
RECALL VARIANCE: 0.02924692172947626

PM4

AVERAGE ACCURACY 0.9521276595744681
AVERAGE PRECISION 0.9556678661004643
AVERAGE RECALL 0.9674085171501403

FLDM1: (avg. accuracy = 96.5%)

In FDLM1 we apply fisher's linear discriminant analysis on the training data assuming gaussian distribution for both positive and negative classes in the univariate dimension. We have final results as shown below. Here we project and reduce the higher dimensional 30d vector into a univariate one by projecting it onto a vector W which maximises the difference of means and minimises the sum of variance of the two classes.

AVERAGE CONFUSION MATRIX [[234. 2.6]
[10.7 128.7]]
AVERAGE ACCURACY: 0.9646276595744682
AVERAGE PRECISION: 0.989010989010989
AVERAGE RECALL 0.9562729873314263

ACCURACY VARIANCE: 0.011844727963168168
PRECISION VARIANCE: 0.010572777282673413
RECALL VARIANCE: 0.030184522628029408

FLDM2: (avg. accuracy = 96.5%)

We get the same results as in FLDM1 here as the only change we make is column shifting which doesn't change any of the things we are measuring. We have final results as shown below.

```
AVERAGE CONFUSION MATRIX [[234.    2.6]
 [ 10.7 128.7]]
AVERAGE ACCURACY: 0.9646276595744682
AVERAGE PRECISION: 0.989010989010989
AVERAGE RECALL 0.9562729873314263
```

```
ACCURACY VARIANCE: 0.011844727963168168
PRECISION VARIANCE: 0.010572777282673413
RECALL VARIANCE: 0.030184522628029408
```

LR1:

In LR1 we are required to implement logistic regression under various parameters as specified below without normalisation.

For sake of comparison here we are using the standard probability threshold of 0.5 and learning rate of 0.01, this case along with all the other cases are covered in code submitted. We are required to compare for Batch Gradient Descent, Mini-batch and Stochastic Gradient Descent.

In LR, we model the distribution of our y , using the sigmoid function, which has w as a parameter. Through gradient descent, we find the optimal w that minimises our negative log likelihood function.

Batch Gradient Descent (avg. accuracy = 57.6%)

We take the average of the gradients of all the training examples and then use that mean gradient to update our parameters. So we just take one step of gradient descent per epoch. The results are as shown below.

LR1

Accuracy

Average = 0.5763440860215054 Standard Deviation = 0.17821262865050122


Precision

Average = 0.7458907294391166 Standard Deviation = 0.31128216099489203

Recall

Average = 0.5396602967368972 Standard Deviation = 0.39119858661610574

ITERATION : 1



```
CONFUSION MATRIX
[[118.  0.]
 [ 58. 10.]]
ACCURACY = 0.6881720430107527
PRECISION = 1.0
RECALL = 0.14705882352941177
```

Mini-Batch Gradient Descent (avg. accuracy = 86.2%)

A mini-batch is a subset of the training data used in each iteration of the training algorithm in mini-batch gradient descent. The results are as shown below.

LR1

Accuracy

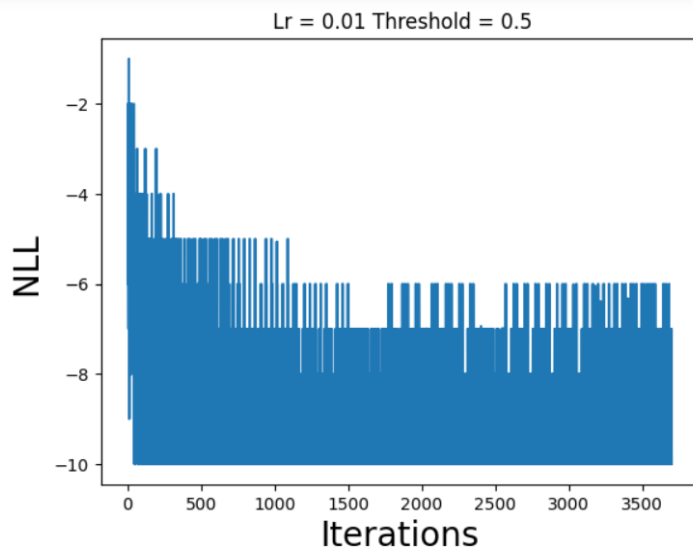
Average = 0.8623655913978494 Standard Deviation = 0.04473686895207466

Precision

Average = 0.7908524555507701 Standard Deviation = 0.12324952986857828

Recall

Average = 0.9003454932901848 Standard Deviation = 0.072502037008662



CONFUSION MATRIX

```
[[89. 29.]
 [ 2. 66.]]
ACCURACY = 0.8333333333333334
PRECISION = 0.6947368421052632
RECALL = 0.9705882352941176
```

Stochastic Gradient Descent (avg. accuracy = 83.4%)

Sequentially modifies the parameters of each training sample in each training sample of the dataset. The results are as shown below.

LR1

Accuracy

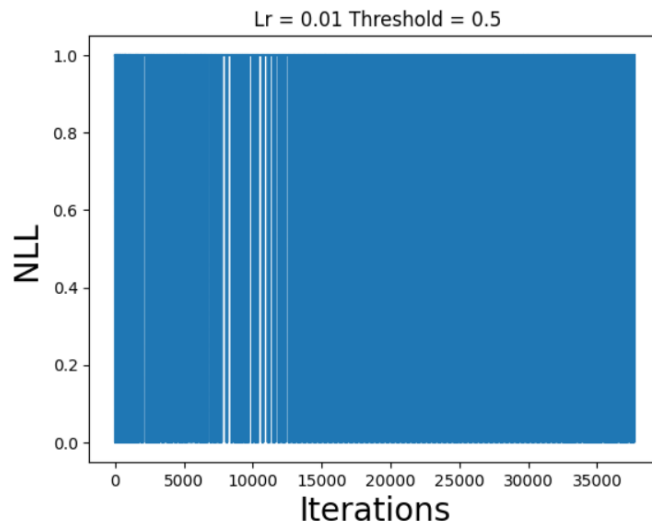
Average = 0.8349462365591398 Standard Deviation = 0.06460790962689322

Precision

Average = 0.7774850747464171 Standard Deviation = 0.16188651462324855

Recall

Average = 0.8794254190447062 Standard Deviation = 0.13402404380875763



CONFUSION MATRIX

```
[[112.  6.]  
 [ 12. 56.]]  
ACCURACY = 0.9032258064516129  
PRECISION = 0.9032258064516129  
RECALL = 0.8235294117647058
```

LR2:

In LR2 we are required to implement logistic regression under various parameters as specified below with normalisation.

For sake of comparison here we are using the standard probability threshold of 0.5 and learning rate of 0.01, this case along with all the other cases are covered in code submitted. We are required to compare for Batch Gradient Descent, Mini-batch and Stochastic Gradient Descent.

Batch Gradient Descent (avg. accuracy = 95.9%)

We take the average of the gradients of all the training examples and then use that mean gradient to update our parameters. So we just take one step of gradient descent per epoch. The results are as shown below.

LR2

Accuracy


Average = 0.9590425531914895 Standard Deviation = 0.012372024840013833

Precision

Average = 0.9487424772768062 Standard Deviation = 0.013953432487262272

Recall

Average = 0.9386084551602998 Standard Deviation = 0.026615303607714334



```
CONFUSION MATRIX
[[116.  4.]
 [  3. 65.]]
ACCURACY = 0.9627659574468085
PRECISION = 0.9420289855072463
RECALL = 0.9558823529411765
```

Mini-Batch Gradient Descent (avg. accuracy =96.8%)

A mini-batch is a subset of the training data used in each iteration of the training algorithm in mini-batch gradient descent. The results are as shown below.

LR2

Accuracy


Average = 0.9686170212765959 Standard Deviation = 0.008393475445776323

Precision

Average = 0.992274523539737 Standard Deviation = 0.007780490700404329

Recall

Average = 0.9201412900703569 Standard Deviation = 0.02607976198504841



CONFUSION MATRIX

```
[[120.  0.]  
 [  6. 62.]]  
ACCURACY = 0.9680851063829787  
PRECISION = 1.0  
RECALL = 0.9117647058823529
```

Stochastic Gradient Descent (avg. accuracy = 93.9%)

Sequentially modifies the parameters of each training sample in each training sample of the dataset The results are as shown below.

LR2

Accuracy

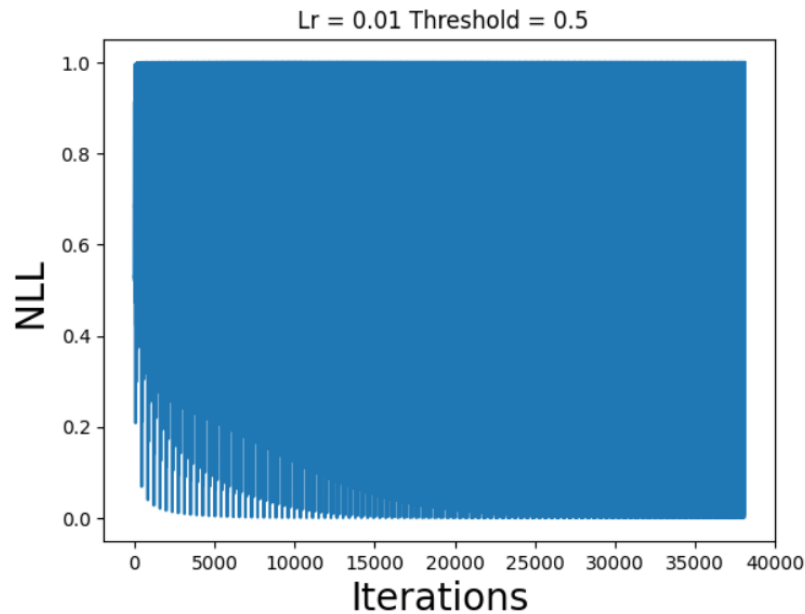
Average = 0.9393617021276597 Standard Deviation = 0.012406280627330424

Precision

Average = 1.0 Standard Deviation = 0.0

Recall

Average = 0.8325692731556942 Standard Deviation = 0.039517369058554366



CONFUSION MATRIX

```
[[120.  0.]  
 [ 11. 57.]]  
ACCURACY = 0.9414893617021277  
PRECISION = 1.0  
RECALL = 0.8382352941176471
```

From the above observations it is clear that LR2 Mini-Batch Gradient Descent is the best performing model with the highest average accuracy at over 96.8%. This is intuitive considering both the feature engineering tasks are performed on data in this model therefore increasing accuracy and also due to the property of how Mini-Batch Gradient Descent is calculated it has a more stable convergence towards global minima since average gradient is calculated over n samples resulting in lower noise (Standard Deviation is the least out of all examples).